

Applied Algorithms: Finding a correlated pair

Martin Aumüller (maau@itu.dk)

Nov 2, 2020

Course Overview

This week:

- Largeish project that incorporates most of what you've learned
- Correct and reasonably fast will probably be a challenge this week

Finding Correlated Elements

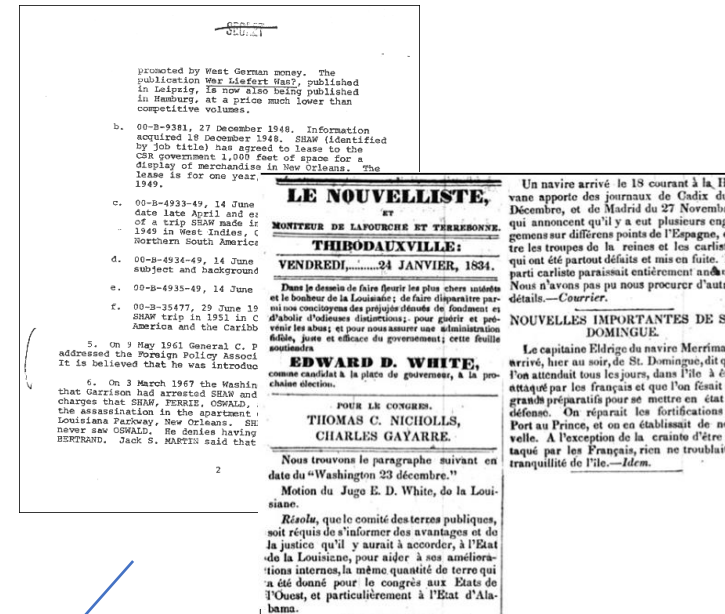

Fairly fundamental and widely used in modern computer science

Plan for today

- Motivate the problem and talk about what it means in general
- Restrict to the specific case we'll be dealing with
- Solutions using algorithms we know
- New solution

Features

Map large images, documents, etc. to digestible vectors



15	0	.2	-3	4	22	1	1
----	---	----	----	---	----	---	---

Example: Similarity Search on Words

- GloVe: learning algorithm to find vector representations for words
- *GloVe.twitter* dataset: **1.2M words**, vectors trained from **2B tweets**, **100 dimension real-valued vector associated with every word**
- Semantically similar words: nearest neighbor search on these vectors



<https://nlp.stanford.edu/projects/glove/>

GloVe Examples (5-NN search for a query)

“denmark”

- “sweden”
- “germany”
- “netherlands”
- “italy”
- “norway”

“copenhagen”

- “brussels”
- “helsinki”
- “belgium”
- “vienna”
- “turin”

“københavn”

- “haugesund”
- “århus”
- “sandefjord”
- “kommune”
- “aalborg”

Features

Goals:

- Easier to compare and classify (can usually search features faster than we can search images)
- Want similar things to stay similar



1	0	1	1	0	0	1	1
---	---	---	---	---	---	---	---



1	0	1	0	1	0	1	1
---	---	---	---	---	---	---	---

Features

Goals:

- Easier to compare and classify (can usually search features faster than we can search images)
- Want different things to stay different



1	0	1	1	0	0	1	1
---	---	---	---	---	---	---	---



0	0	0	0	1	0	1	0
---	---	---	---	---	---	---	---

Using Features

- We want to perform operations on these vectors
- Similarity queries (reverse image search), find close pair (plagiarism detection), much more complicated operations

Bit Vectors

- Again: inputs are bit vectors
- Similar bit vectors = high *inner product*
 - A.k.a. the number of 1's in common

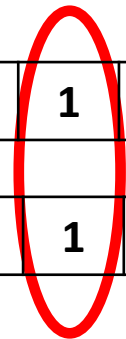
1	0	1	1	0	0	1	1
1	0	1	0	1	0	1	1

Recall: Inner Product =
number of ones in x & y

Cheap!

Bit Vectors

- Dissimilar bit vectors = low *inner product*
 - A.k.a. the number of 1's in common



1	0	1	1	0	0	1	1
0	0	0	0	1	0	1	0

Our Problem Today

- Given a list of bit vectors, and a given similarity (inner product), find the pair with at least that similarity.
 - Similarity is given to you (will always be the same in all tests)
 - Guaranteed to have one unique pair

Problem kind of orthogonal to “Orthogonal Vectors”

Our Problem Today

- Given a list of bit vectors, and a given similarity, find the pair with at least that similarity.

Find pair with similarity ≥ 3

1	0	1	0	1	0	1	1
---	---	---	---	---	---	---	---

0	0	0	0	1	0	1	0
---	---	---	---	---	---	---	---

1	0	1	1	0	0	1	1
---	---	---	---	---	---	---	---

0	0	1	1	0	1	0	0
---	---	---	---	---	---	---	---

Our Problem Today

- Given a list of bit vectors, and a given similarity, find the pair with at least that similarity.

1	0	1	0	1	0	1	1
---	---	---	---	---	---	---	---


0	0	0	0	1	0	1	0
---	---	---	---	---	---	---	---

1	0	1	1	0	0	1	1
---	---	---	---	---	---	---	---

0	0	1	1	0	1	0	0
---	---	---	---	---	---	---	---

Find pair with similarity ≥ 3

Similarity 4



Our Problem Today

- Given a list of bit vectors, and a given similarity, find the pair with at least that similarity.

Find pair with similarity ≥ 3


1	0	1	0	1	0	1	1
---	---	---	---	---	---	---	---

0	0	0	0	1	0	1	0
---	---	---	---	---	---	---	---

1	0	1	1	0	0	1	1
---	---	---	---	---	---	---	---

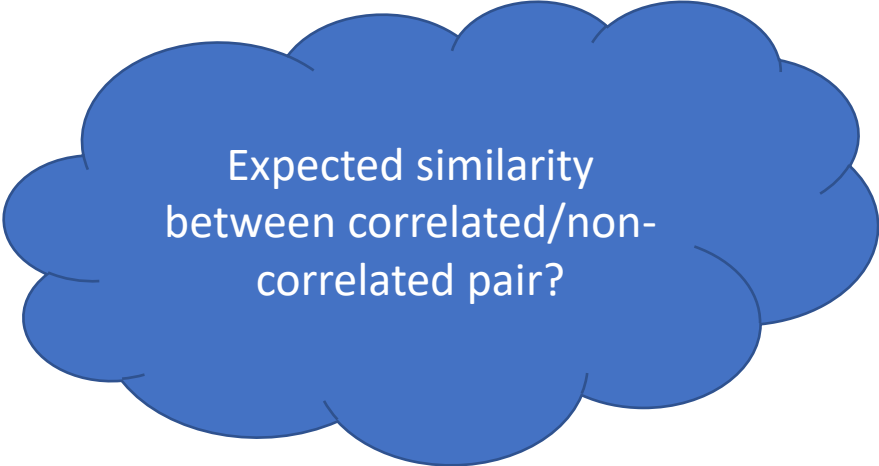
0	0	1	1	0	1	0	0
---	---	---	---	---	---	---	---

Similarity 1



Simplifying The Problem

- Similarity threshold is given
- Only one close pair
- All vectors have same length ℓ
- All bit vectors are generated randomly:
 - $n-1$ of them are generated at random; each bit is independently set with probability $1/3$ (expect $\ell/3$ ones)
 - In particular: each bit is set with the same probability
 - Then we generate one similar to a previous vector. With probability $7/8$ it has the same bit; otherwise it's random.



Expected similarity
between correlated/non-
correlated pair?

Simplifying The Problem

- Similarity threshold is given (≥ 70)
- Only one close pair
- All vectors have same length (256 bits)
- All bit vectors are generated randomly

These assumptions should hopefully make the code easier to write

How do we solve this?

Simple brute force algorithm?

- Compare each pair of vectors
- $O(n^2)$ comparisons
- How do we compare two 256-dim bit vectors quickly?
 - 4 bitwise ANDs
- Can evaluate $n = 100\,000$ vectors in about 35 seconds on my laptop
- Let's do this together!

How do we improve over this?

Heuristics

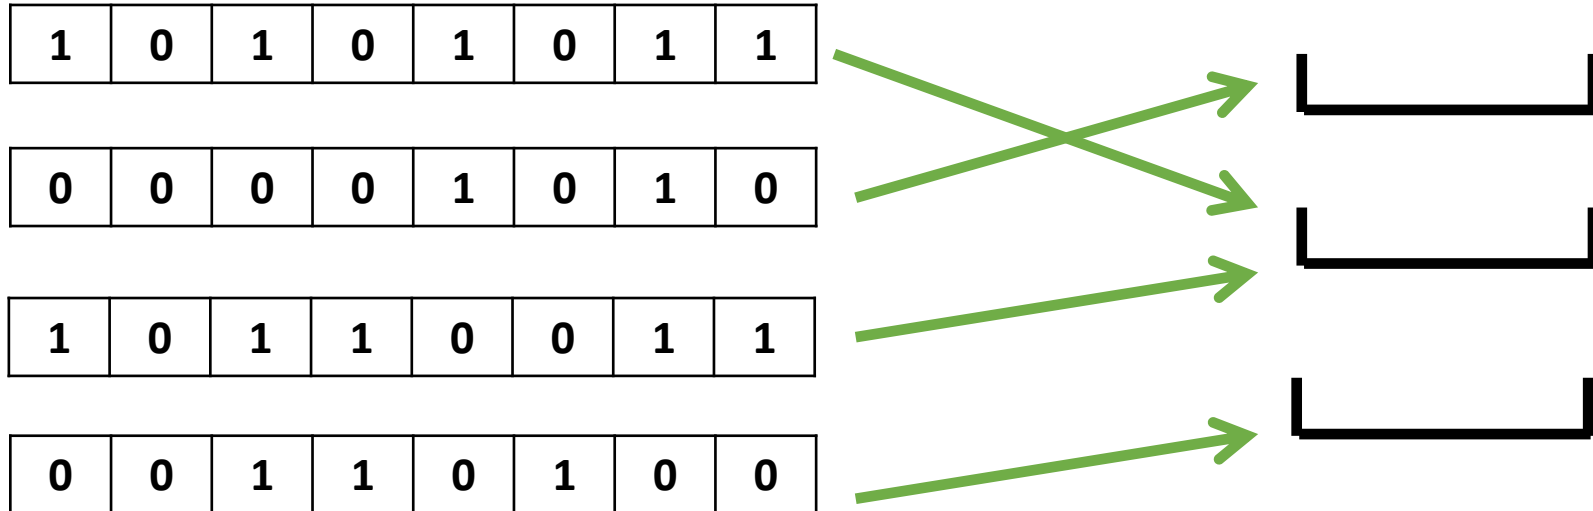
- Be creative
- I tried stopping each comparison short if the first 64 bits (the first AND) had low correlation. Gave ~50% speedup
 - Only OK because the data is generated randomly. Could make us miss the close pair on certain datasets

MinHash

- Way to hash bit vectors while retaining information about their similarity
 - “Locality-sensitive hashing”
- Used fairly widely in practice
- We will modify it a little bit today for our purposes


Goal of MinHash

- Hash vectors to buckets
- Only compare two vectors if they are in the same bucket
- Notion: Two vectors “*collide*” if they are in the same bucket



Goal of MinHash

- Hash vectors to buckets
- Only compare two vectors if they are in the same bucket
- Notion: Two vectors “*collide*” if they are in the same bucket



Only need to compare these two

Goal of MinHash

- Hash vectors to buckets
- Only compare two vectors if they are in the same bucket
- Different to HyperLogLog hashing:
 - Want that similar items collide!
 - Not spread them out randomly

What we need:


- Buckets shouldn't be too big
- Similar pair should be likely to be in the same bucket.

If we don't find the similar pair, try again with new hash functions

MinHash

- Choose your MinHash function: A random permutation of $\{1, \dots, 256\}$
- To hash: Take the position of the first 1 in the vector encountered by traversing the permutation
- That position is the MinHash of the vector

Random
permutation:

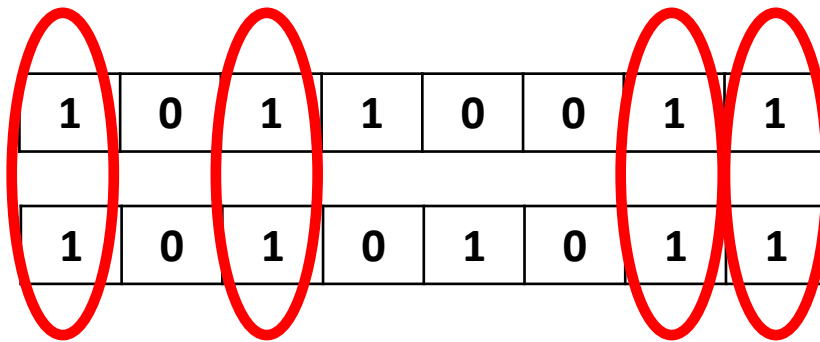


MinHash of this vector is 5,
because the first 1 we found
is in position 5

Very brief analysis

- Probability that x and y hash together is proportional to $\text{similarity}(x,y)$
- The precise value is the Jaccard Similarity:

$(\# \text{ 1s in common}) / (\text{total positions that have a 1 in either})$



1	0	1	1	0	0	1	1
1	0	1	0	1	0	1	1

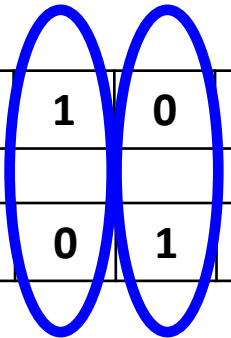
The diagram shows a 2x8 grid of binary digits. Red ovals are drawn around the columns where both the top and bottom rows have a '1'. These columns are at indices 0, 2, 6, and 7 (0-indexed from the left). The first oval encloses the first column (1,1), the second oval encloses the third column (1,1), the third oval encloses the seventh column (1,1), and the fourth oval encloses the eighth column (1,1).

If one of these is the min
they collide

Very brief analysis

- Probability that x and y hash together is proportional to $\text{similarity}(x,y)$
- The precise value is the Jaccard Similarity:

$(\# \text{ 1s in common}) / (\text{total positions that have a 1 in either})$



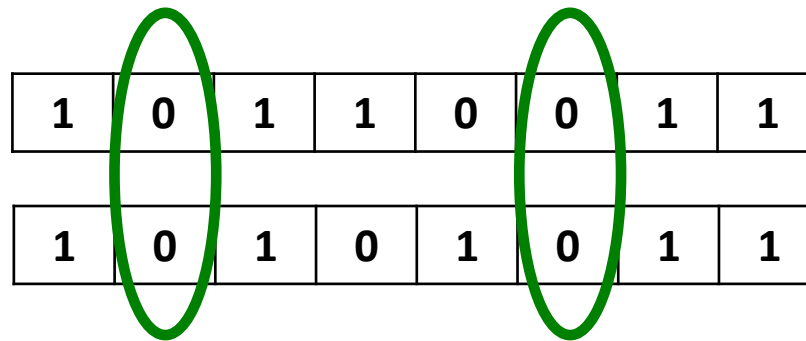
1	0	1	1	0	0	1	1
1	0	1	0	1	0	1	1

If one of these is the min
they don't collide

Very brief analysis

- Probability that x and y hash together is proportional to $\text{similarity}(x,y)$
- The precise value is the Jaccard Similarity:

$(\# \text{ 1s in common}) / (\text{total positions that have a 1 in either})$



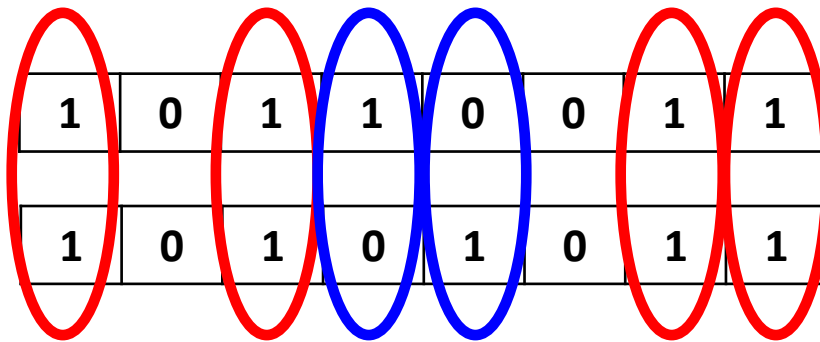
1	0	1	1	0	0	1	1
1	0	1	0	1	0	1	1

These can never be the min
so we ignore them

Very brief analysis

- Probability that x and y hash together is proportional to $\text{similarity}(x,y)$
- The precise value is the Jaccard Similarity:

$(\# \text{ 1s in common}) / (\text{total positions that have a 1 in either})$



1	0	1	1	0	0	1	1
1	0	1	0	1	0	1	1

6 things to choose from, 4 collisions

$$\text{Prob}(\text{collision}) = 4/6 = 2/3$$

How much does this help us?

- How likely are we to see two random vectors collide?
- Order of magnitude when hashing 1 million elements:
 - Bucket size $\sim 100k$
 - Bucket size $\sim 1k$
 - Bucket size ~ 1
- Pretty likely to collide—something like probability $1/3$ (so bucket size $\sim 100k$)
- This isn't worth it, so what do we do?

Hash several times

- Use new hash function to split our buckets further (in other words, a new random value for each bit *position*)
 - Have to be careful—have to choose new hash functions, not reuse old ones!
- Two vectors only collide if they collide in each of our several hashes

Hash several times

Hash = 5

3	0	5	1	7	4	6	2
---	---	---	---	---	---	---	---

0	0	1	0	1	1	1	1
---	---	---	---	---	---	---	---

Hash = 6

6	2	0	7	1	4	5	3
---	---	---	---	---	---	---	---

0	0	1	0	1	1	1	1
---	---	---	---	---	---	---	---

Hash = 7

0	7	1	3	2	4	5	6
---	---	---	---	---	---	---	---

0	0	1	0	1	1	1	1
---	---	---	---	---	---	---	---

So our MinHash is (5,6,7). We will only compare to vectors that also hash to (5,6,7)

Hash several times

- Our collision probability is getting low. What happens when we compare all to all and we don't find the pair of vectors?
- Repeat again with a new hash
- If we know the close vectors have probability p of collision, how many expected repetitions do we need?
 - $1/p$
- Don't need to parameterize by this; can just keep trying until you get the vector

How big do we want our buckets?

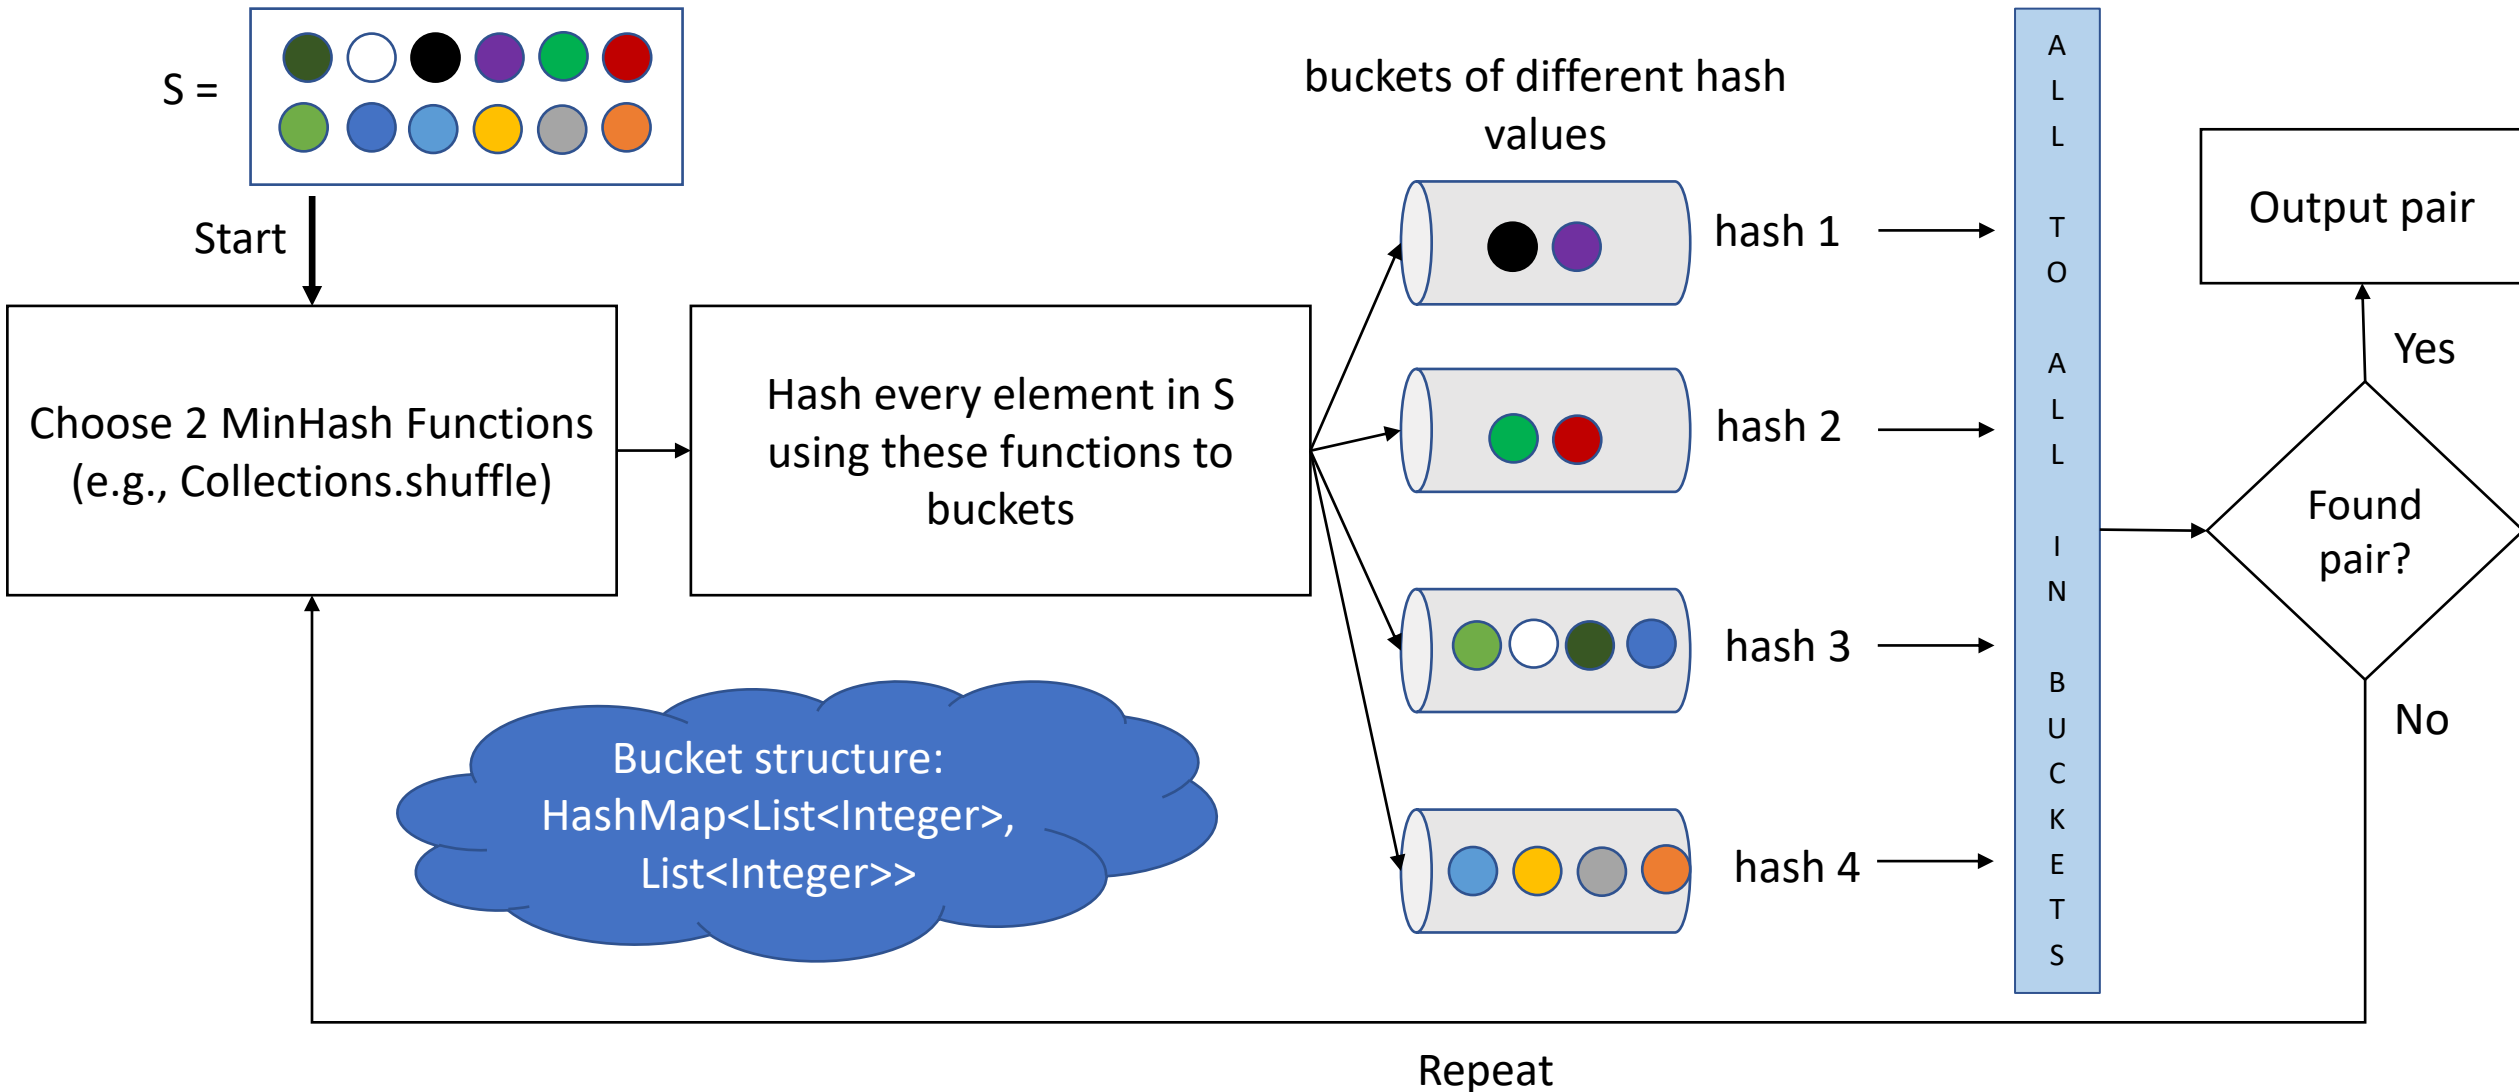
- Definitely small enough to fit in memory
- Run experiments to see what is fastest
- How to control bucket size?
 - Different # concatenated hashes



Implementation

- All-compare-all within a bucket uses the naïve method which we have already implemented
 - Just need to get the vectors into an array (or vector, etc)
- How do you allocate buckets?
 - Dynamically (?)
 - Can count bucket sizes first
- How do you *label* buckets?
 - Don't want to create a bucket for all possible MinHashes (most will be empty)
 - Use hash table?
 - Only 256 positions, i.e., MinHash has 8 bits. Store many of them as one 32-bit or 64-bit integer?

Overview w/ 2 MinHash functions on input S



Speeding Up

- MinHash is much much faster can do 50 000 vectors in 0.3 seconds; naïve takes 184 seconds
- This speedup is probably language-independent, but need to make sure you don't lose time

Assignment

- Input: List of 256-dimensional vectors, each listed as 4 signed 64-bit integers, from stdin, randomized as described before
- Output: index of the correlated pair, smaller index first
- CodeJudge
 - No race
 - Largest instance you can solve in 10 seconds
 - 1k, 10k, 100k, etc
 - Should be fairly tamper-proof; let me know if it's not